

Plugin Writer's Guide	
Author	Nick Battle
Date	04/10/22
Issue	0.1

0 Document Control

0.1 Table of Contents

0	Document Control.....	2
0.1	Table of Contents.....	2
0.2	References.....	2
0.3	Document History.....	3
0.4	Copyright.....	3
1	Overview.....	4
1.1	VDMJ.....	4
1.2	VDMJ Interfaces.....	4
1.3	LSP Server Plugin Architecture.....	4
2	Server/Plugin Interactions.....	6
2.1	Server Initialization.....	6
2.2	Parsing and Checking.....	6
2.3	Opening a File.....	7
3	Analysis Plugin Implementation.....	8
3.1	The Build Environment.....	8
3.2	Plugin Configuration.....	8
3.3	Plugin Class Construction.....	8
3.4	The Plugin Registry.....	9
3.5	Event Handling.....	9
3.6	Client Capabilities.....	11
3.7	Server Startup.....	11
3.8	Other Plugin Processing.....	11
4	Example Plugin Functionality.....	13
A.	Events.....	14

0.2 References

- [1] Wikipedia entry for The Vienna Development Method, http://en.wikipedia.org/wiki/Vienna_Development_Method
- [2] Wikipedia entry for Specification Languages, http://en.wikipedia.org/wiki/Specification_language
- [3] VDMJ, <https://github.com/nickbattle/vdmj>
- [4] <https://microsoft.github.io/language-server-protocol/overviews/lsp/overview/> and <https://microsoft.github.io/debug-adapter-protocol/overview>
- [5] Visual Studio Code, <https://github.com/microsoft/vscode>
- [6] VDM VSCode extension, <https://github.com/overturetool/vdm-vscode>

0.3 Document History

Issue 0.1 04/10/22 First draft.

0.4 Copyright

Copyright © Aarhus University, 2022.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.3 or any later version published by the Free Software Foundation; with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts.

1 Overview

This document describes how to write *Analysis Plugins* for the VDMJ LSP language server.

Section 1 gives an overview of the architecture into which plugins fit. Section 2 walks through various common scenarios to describe the interaction of plugins with the server. Section 3 gives detailed information about how to implement plugins. Section 4 gives some examples of what would be possible with plugins and how to achieve it. Lastly, Section 5 describes how the language server is likely to change in future releases.

1.1 VDMJ

VDMJ provides basic tool support for the VDM-SL, VDM++ and VDM-RT specification languages, written in Java [1][2][3]. It includes a parser, a type checker, an interpreter and debugger with coverage recording, a proof obligation generator, user definable annotations and a combinatorial test generator, as well as *JUnit* support for automatic testing.

1.2 VDMJ Interfaces

VDMJ offers language services independently of the means used to access those services. That means that VDMJ can be used from several different IDE environments.

A basic user interface is built in to VDMJ and offers a simple command line. But VDMJ can also be used via the LSP/DAP protocols [4] and so can be used by an IDE like Visual Studio Code [5][6]. To support this, VDMJ includes a “language server” that responds to LSP/DAP connections.

1.3 LSP Server Plugin Architecture

Services are added to the language server via a number of *Analysis Plugins*, which are responsible for all processing that relates to a particular *analysis*. An analysis is an independent aspect of the processing of a VDM specification. For example, the fundamental analyses cover parsing, type checking and interpretation. But an analysis could also be a translation from VDM to another language, or a more advanced kind of type checking or testing.

Plugins for the fundamental analyses are built into the language server, but additional plugins can be written independently and added to the language server environment easily. This document describes how to write such plugins.

The basic architecture of the VDMJ language server is shown in Figure 1. An IDE “Client” connects to a JSON RPC server, which uses a *Workspace Manager* to handle RPC calls, one for LSP, one for DAP. The *Workspace Manager* constructs a number of plugins at startup, some of which are built-in and some of which are user configured. All plugins register with a *PluginRegistry* and with an *EventHub*. The registry allows *Workspace Managers* to make a request from all plugins – for example, the *getCodeLenses* call illustrated in the Figure would collect all of the lenses that the configured plugins define. The *EventHub* allows the *Workspace Managers* to inform plugins about activity from the Client, where individual plugins choose to subscribe to particular event types.

Plugins make use of VDMJ language services to actually do their processing (parsing, type checking etc). The *Workspace Managers* should not directly make calls to VDMJ¹.

¹ though at release 4.5.0 there are still some direct calls that have yet to be migrated to the appropriate plugin.

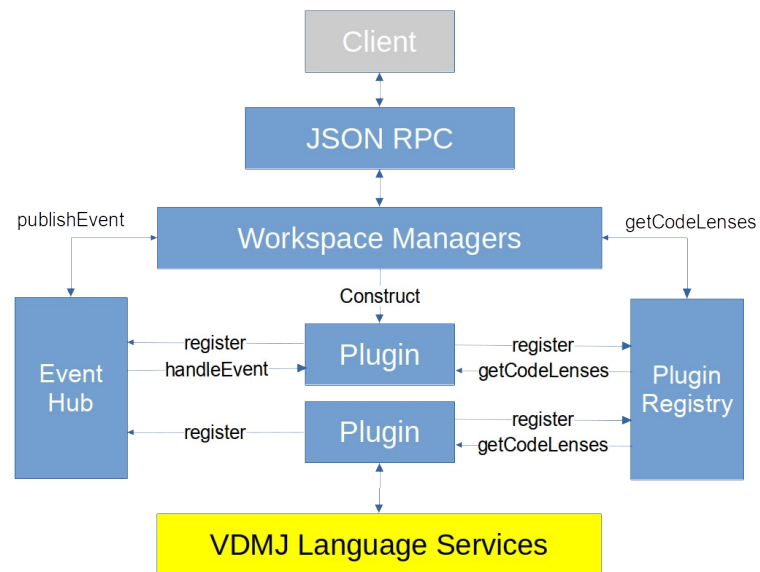


Figure 1: The VDMJ Language Server Architecture

2 Server/Plugin Interactions

This section shows the sequence of manager/plugin interactions for a number of common scenarios, in the hope that this will clarify how plugins provide their services. A more detailed description of how to write plugins in Java is given in the next section.

2.1 Server Initialization

The language server is typically initialized when the IDE opens a VDM project and perhaps opens the first file. The LSP protocol [4] defines a number of interactions at initialization, which causes the following:

1. The JSON RPC server starts and listens for connections.
2. The IDE sends the LSP “initialize” request, which is sent to the *LSPWorkspaceManager*, passing the Client’s capabilities and the root location of the project.
3. The Workspace Manager’s construction registers the built-in plugins with the *PluginRegistry*. This calls their *init* methods, which may register the plugins to receive various event types with the *EventHub*.
4. The Client capabilities and project root are stored in the Workspace Manager.
5. The *LSPXWorkspaceManager* is called to register user-defined plugins with the *PluginRegistry*, and they may also register with the *EventHub*.
6. The project files are (recursively) loaded into Workspace Manager memory. This includes the generation of any external format files, and respects the *vdmignore* list or *ordering* files.
7. The server capabilities for the built-in plugins is prepared and then the *setServerCapabilities* method of each plugin is called, via the *PluginRegistry*, to allow them to add or change settings.
8. An *InitializeEvent* is published, and any responses collected.
9. The final server capabilities and any responses to the event are returned to the IDE.
10. The IDE should respond with an “initialized” LSP message, which is dispatched to the Workspace Manager.
11. The memory-loaded project files are parsed and type checked. See 2.2 .
12. An *InitializedEvent* is published, and any responses collected. This can be used by plugins to make dynamic registrations for LSP features that they offer.
13. Any warning or error messages from the parse/check, along with any responses to the event are returned to the IDE.
14. The system is then quiet until the user does something in the IDE.

2.2 Parsing and Checking

At several points while interacting with the IDE, the currently loaded specification must be re-parsed and type checked, sending any warnings or errors to the IDE as notifications. The process is event driven, as follows:

1. Some other processing calls the *checkLoadedFiles* method of the *LSPWorkspaceManager*.
2. A *CheckPrepareEvent* is published, which informs plugins that a parse/check is about to happen. They can use this to clear their state, in anticipation of more events.
3. A *CheckSyntaxEvent* is published and the responses collected. Plugins are expected to parse the specification. If there are errors in the responses, the check is abandoned at this point and the collected messages returned to the IDE. If there are no messages or only warnings, processing continues.

4. A *CheckTypeEvent* is published and the responses collected. Plugins are expected to type check the specification. If there are errors in the responses, the check is abandoned at this point and the collected messages returned to the IDE. If there are no messages or only warnings, processing continues.
5. A *CheckCompleteEvent* is published and the responses collected. Plugins are expected to initialize any runtime components that they have (e.g. the interpreter resets itself).
6. Any responses collected so far are returned to the IDE as notifications.
7. A "slsp/checked" notification is also sent to the IDE.

Note that events are processed by plugins in order, with the built-in plugins coming first. So user-defined plugins can assume that (for example) the specification has been parsed by the AST plugin before it handles the *CheckSyntaxEvent*. The ordering of plugins is discussed more in section 3.2.

2.3 Opening a File

It may be that the server is initialized on the opening of the first source file in the IDE, in which case the processing above occurs as beforehand, but the following events occur when an existing source file is opened in the IDE:

1. A "textDocument/didOpen" LSP notification is sent to the language server, which is handled by the *LSPWorkspaceManager*. Notifications do not expect responses.
2. The file is checked against any *vdmignore* list.
3. The Workspace Manager notes that the file is open and whether it is a new file (whether it was loaded at startup, or a newly created file).
4. An *OpenFileEvent* is published via the *EventHub*. The responses are ignored, because the open file LSP message is a notification.
5. If the file is new, the project is re-parsed and checked, and any warnings or errors sent to the Client. See 2.2.
6. The built-in plugins will have registered a code lens capability with the Client, which will send a "textDocument/codeLens" request, if it also supports lenses. This is handled by the *LSPWorkspaceManager*.
7. Responses to a code lens request are generated by calling the *PluginRegistry* to ask each plugin whether it wants to contribute lenses for the open file. The AST and TC plugins will contribute their "Launch | Debug" lenses, but there may be others.
8. The system is then quiet until the user does something else in the IDE.

2.4 Editor Changes are Made

When the user updates a file, even if those changes have not been saved, the "edits" are sent to the language server so that it can keep its in-memory version of the file up to date. The following occurs:

1. When you update a file, the Client sends "textDocument/didChange" notifications to the language server, which are handled by the *LSPWorkspaceManager*.
2. The Workspace Manager updates the in-memory version of the file, respecting the *vdmignore* list (so you can edit any file, but only VDM changes get processed further).
3. If you are editing an external format generated file, an IDE warning notification is sent.
4. A *ChangeFileEvent* is published. This can be handled by any plugin, but the AST plugin will use it to re-parse the updated in-memory file and send back syntax errors and warnings, so that these appear on the fly.

2.5 A File is Created, Deleted or Saved

Edits made in an open editor are sent to the language server, as discussed above, but this is not the same as the file changes being made permanently on disk. When the user saves the changes, or a new file is created or deleted, the following events occur:

1. The built-in plugins register for “changedWatchedFiles” LSP events at startup. The Client will send these whenever filesystem changes are noticed within the project.
2. A file save or new file create/delete sends a “workspace/didChangeWatchedFiles” LSP notification. The notification is either of type “create”, “change” or “delete” and includes folders as well as files.
3. The server goes through all of the changes received and, depending on what changes, decides whether afterwards it needs to do nothing, to just re-check the in-memory specification or completely reload and re-check. See 2.2 .

Note that there are currently no events published for watched file changes. This may be added in a future release.

2.6 An Interpreter Session is Opened

Execution and debugging of a specification uses the DAP protocol rather than the LSP protocol, and there is a *DAPWorkspaceManager* that handles these messages. But the basic architecture of the DAP interactions is the same as for LSP. When an executable session is opened, the following occurs:

1. A DAP “initialize” request is sent to the language server. As with LSP, this include the Client’s debugging capabilities and these are saved by the *DAPWorkspaceManager*.
2. A fixed set of server capabilities are sent back to the Client. Unlike LSP, the plugins cannot add or change capabilities currently. This may change in a future release to allow plugins to offer advanced debugging features.
3. A *DAPInitializeEvent* is published, and any responses saved.
4. An “initialized” DAP event is returned to the Client, along with any responses to the event above.
5. The Client may then send DAP requests to set breakpoints, if these have been set in the IDE. These are passed to the VDMJ interpreter.
6. The Client will then send a DAP “launch” command, which indicates whether the session is for debugging or just execution, along with other settings. These are saved by the Workspace Manager.
7. The Client should then send a “configurationDone” DAP notification. This is the indication that everything is set up and VDMJ can initialize the specification. Before this is done, a *DAPConfigDoneEvent* is published.
8. The “IN” plugin is obtained from the PluginRegistry, and this provides a VDMJ *Interpreter* instance. The specification is then initialized using a background thread.
9. The initialization thread will send DAP events to print the VDM banner on the debug console, and then run the *init* method of the *Interpreter*. When the initialization completes, the duration of the initialization is printed and the language server waits for the user to type something on the console.
10. (If a command was sent with the “launch” request, this will be executed immediately, rather than waiting for the user).
11. When the user enters a line on the console, an “evaluate” request is sent to the language server. The line typed is passed to all plugins, via the *PluginRegistry*, to call their *getCommand* methods. If a plugin recognises the line, it returns a *Command* object. If there

are usage errors, it can return an *ErrorCommand* or just throw an exception.

12. A *DAPEvaluateEvent* is published, with details of the *Command* being executed.
13. The *Command's run* method is invoked by the Workspace Manager. This is responsible for doing whatever the command intends, and may use other plugins or the Workspace Manager to achieve this.
14. The responses from the *Command run* are sent back to the Client, which typically displays them on the console.

Note that the DAP protocol is less well integrated with the plugin architecture, with fewer events being published for example. This will be improved in future to allow tighter integration of plugins with executions.

2.7 Unrecognised LSP/DAP Messages

It is possible that a new IDE feature wants to send a new message to the language server, which is not supported by the LSP or DAP protocols. This already happens with the *VDM VSCode* extension [6], where several non-standard messages are sent via LSP.

Typically, new messages will be handled by new plugins in the language server. So to connect the two, the *LSPXWorkspaceManager* and *DAPXWorkspaceManager* publish events, *UnknownMethodEvent* and *UnknownCommandEvent* for LSP and DAP respectively². Plugins can register for these events and check for the new message/command names that they support.

² LSP calls its RPC names “methods”, while DAP calls them “commands”.

3 Analysis Plugin Implementation

This section describes how to implement a new plugin for the VDMJ language server. The description is based around an example plugin that is provided with VDMJ.

3.1 The Build Environment

VDMJ is written in Java, so it expects plugins to be written in Java also. You can choose the version of Java that you use, but it must be at least version 8 as the language server itself requires this.

The example plugin provided with VDMJ is built using Maven, which defines the dependencies required. But you may prefer to write your plugin using a different dependency management system.

The VDMJ source is available here [3]. The example plugin is in *examples/lspplugin*. The entire suite can be compiled with the Maven command “mvn clean install”.

3.2 Plugin Configuration

All analysis plugins extend the abstract Java class *workspace.plugins.AnalysisPlugin*.

To configure the language server with the user supplied plugins to be used, the Java property *lsp.plugins* must be set to a comma-separated list of the class names of the plugins. The fundamental plugins are configured automatically, and *ahead* of all user supplied plugins.

The order of plugins, which may be significant (see below), is as follows:

1. AST plugin (parser)
2. TC plugin (type checker)
3. IN plugin (interpreter)
4. PO plugin (proof obligations)
5. CT plugin (combinatorial testing)
6. User supplied plugins in order of the *lsp.plugins* property.

To include just the example plugin, the property would be set as follows (using -D for the java command line):

```
-Dlsp.plugins=examples.ExamplePlugin
```

Naturally, as well as being listed in *lsp.plugins*, the classes must also be available on the classpath. Typically, plugins are compiled into separate jars which are added to the classpath, for example (again, using the java command line).

```
-cp <etc>:examples/lspplugin/target/lspplugin-4.5.0-SNAPSHOT-221001.jar
```

The way of setting *lsp.plugins* and the classpath differs between IDEs. In VDM VSCode, they are set via the “Settings” page for the extension (which edits the *.vscode/settings.json* file).

3.3 Plugin Class Construction

When the language server starts, a single instance of each plugin is constructed. The construction of plugins looks for a method in each plugin class with this signature:

```
public static <plugin class> factory(Dialect dialect)
```

If that method exists, it is passed the *dialect* of the language server, which allows the plugin to create a different variant for each dialect – remember VDM-SL specifications include modules, and other

dialects include classes, so there are significant differences. Dialect subclasses can be more efficient than testing the server dialect repeatedly. If this factory method does not exist, the default class constructor is used.

The example plugin implements the factory method, returning one of two subclasses:

```
public static ExamplePlugin factory(Dialect dialect)
{
    switch (dialect)
    {
        case VDM_SL:
            return new ExamplePluginSL();

        case VDM_PP:
        case VDM_RT:
            return new ExamplePluginPR();

        default:
            Diag.error("Unsupported dialect " + dialect);
            throw new IllegalArgumentException(...);
    }
}
```

3.4 The Plugin Registry

Having constructed a plugin, the *registerPlugin* method of the *PluginRegistry* is called to register the plugin. It does the following:

```
public void registerPlugin(AnalysisPlugin plugin)
{
    plugins.put(plugin.getName(), plugin);
    plugin.init();
    Diag.config("Registered analysis plugin: %s", plugin.getName());
}
```

Note that this requires two methods to be implemented: *getName* and *init*.

The *String* name of your plugin can be used by other plugins to obtain services and data that you may provide. By convention, it is a short string that is also reflected in the class names of the plugin, but this is not a requirement. For example, the name of the parser plugin is “AST”, the type checker plugin is “TC” and the interpreter is “IN”. Plugin names must be unique in a running server, though note that a user plugin could replace a built-in plugin with the same name.

The *init* method, as the name suggests, should initialize your plugin. Typically, this will involve registering for various *events* with the *EventHub* (see 3.5) and setting local fields.

After registration, your initialized plugin will be invoked when various things happen in the language server, either by events that it registered to receive, or by methods being called as discussed below.

Plugin instances can subsequently be found by calling the *getPlugin* method of the registry.

3.5 Event Handling

One way for the language server to communicate with plugins is via *Events*. These are published by the language server when various protocol events occur; plugins subscribe to particular events of interest. An *EventHub* manages the plugin subscriptions and distributes *Events*.

If they register to receive events, plugins must implement the *EventListener* interface. This defines two overloaded methods called *handleEvent*, which are passed either an *LSPEvent* or a *DAPEvent*. Events indicate that something has happened in the language server and carry information about that event.

For example, the *ASTPlugin* does the following in its *init* method:

```
public abstract class ASTPlugin
    extends AnalysisPlugin implements EventListener
{
    @Override
    public void init()
    {
        eventhub.register(InitializedEvent.class, this);
        eventhub.register(ChangeFileEvent.class, this);
        eventhub.register(CheckPrepareEvent.class, this);
        eventhub.register(CheckSyntaxEvent.class, this);
        this.dirty = false;
    }
}
```

The *eventhub* field is available to all plugins and is the same as *EventHub.getInstance()*. The register method is passed an *Event* subclass to subscribe to, and an *EventListener* to handle those events. The simplest design is for the *EventListener* to be implemented by the plugin itself, but you can use a separate listener if you wish. The listener interface defines two *handleEvent* methods.

Having registered for an *Event*, the appropriate *handleEvent* method is called whenever the event occurs. So for example, *ASTPlugin* has:

```
@Override
public RPCMessageList handleEvent(LSPEvent event) throws Exception
{
    if (event instanceof InitializedEvent)
    {
        return lspDynamicRegistrations();
    }
    else if (event instanceof ChangeFileEvent)
    {
        return didChange((ChangeFileEvent) event);
    }
    else ...
}
```

Notice that *handleEvent* returns an *RPCMessageList*, which is a list of *JSONObject*s that represent the response to the Client, if any (you can return null). The responses from every plugin that is registered for an *Event* are collected and sent back to the Client, along with any standard responses from the language server itself.

The *EventHub* calls each of the registered plugins, in the order of their registration, on the same thread – the main LSP or DAP listening thread. This means that if it takes a long time to process an event, you will be holding up the rest of the language server. In that case, you should consider using a background thread to perform the work. The language server includes a *CancellableThread* class that may help, but that is beyond the scope of this document.

In the example above, the *InitializedEvent* is sent when the language server is exchanging initialize messages with the Client. By handling this event, the *ASTPlugin* can add some dynamic registrations for services that it requires but which cannot be set via the standard initialize response.

The example plugin registers itself to receive every possible event type, and prints out the name of the event when it handles one:

```
@Override
public RPCMessageList handleEvent(LSPEvent event) throws Exception
{
    System.out.println("ExamplePluginSL got " + event);
    return null;
}
```

A full list of Event types and when they are raised is given in Appendix A.

3.6 Client Capabilities

The IDE Client send its LSP capabilities to the language server as part of the “initialize” request. These are cached in the server and can be accessed from plugins via the *LSPWorkspaceManager*:

```
LSPWorkspaceManager manager = LSPWorkspaceManager.getInstance();
```

There are two methods on the manager to look at Client capabilities:

1. *boolean hasClientCapability(String dotName)*. This method takes a dot-format capability name (like “textDocument.synchronization.dynamicRegistration”) and returns a boolean, which is true if the capability exists and is “true”, else false.
2. *<T> T getClientCapability(String dotName)*. This method takes a dot-format capability name and returns the actual value of that capability, or null if it does not exist.

3.7 Server Startup

3.7.1 Setting Server Capabilities

When the language server starts, it has various LSP protocol exchanges with the Client (the IDE). These inform the server of the Client’s capabilities and allow the server to inform the Client of its own capabilities.

Plugins may want to extend the capabilities of the language server, so they need to be able to contribute to this exchange. To enable this, plugins can implement the following method:

```
@Override  
public void setServerCapabilities(JSONObject capabilities)
```

The *JSONObject* passed is the server capabilities response that has been built so far. You will see that it includes capabilities supported by the fundamental plugins, as well as any plugins that are earlier than you in the plugin configuration (see 3.2). This object can be amended by the plugin to set capabilities that it provides itself.

3.7.2 Initialization Events

Two events are sent during startup:

1. *InitializeEvent*. This occurs when the server has received the “initialize” LSP request. All plugins have been initialized, and the Client capabilities are available for query. All of the files in the project will have been discovered and cached – including the reading of external file formats. The responses to the event are sent to the Client along with (and after) the “initialize” response.
2. *InitializedEvent*. This occurs when the server has received the “initialized” LSP request. The loaded files will have been type checked and any errors or warning notifications will be ready to send to the Client. Any responses to the event are sent along with (and after) any type checking errors.

3.8 Other Plugin Processing

3.8.1 Code Lenses

When a file is opened in the IDE, it is possible that the file includes *code lenses*, which are annotations that appear in the Editor as clickable items to perform useful code related activities. Code lenses can be provided by plugins, and when the IDE opens a file, the following method is called on

all plugins (the abstract version in *AnalysisPlugin* does nothing):

```
public JSONArray getCodeLenses(File file)
```

The *JSONArray* returned contains any code lenses that the plugin wishes to create for the file name passed. The *CodeLens* class contains useful methods for helping to create these.

For example, *ASTPlugin* and *TCPlugin* implement this method and return lenses that offer the “Launch | Debug” lenses that appear above executable functions or operations. The *ExamplePlugin* implements this method and adds a “Config” lens for explicit function definitions only. If you click the Config lens, VSCode will open the *launch.json* file in an editor window.

Note that the command sent back to the Client with a lens is IDE specific, so lenses would normally test the Client type, with a test like *isClientType("vscode")*, before returning a lens.

3.8.2 Plugin Commands

Plugins can also contribute commands that can be used in the “Debug Console” window (in VSCode) when an executable session is open. To do this, plugins implement this method:

```
public Command getCommand(String line)
```

This is called whenever the user types a line in the console. It is responsible for parsing the line and, if recognised, returning a *Command* subclass that implements the command via its *run* method. If the line is not recognised by the plugin, a null should be returned. If the line is recognised, but malformed, a Java *IllegalArgumentException* can be thrown, or an *ErrorCommand* object can be returned.

If multiple plugins recognise the same line, the last (in order, see 3.2) is used. Note that this means a user plugin can replace a standard command from one of the earlier built-in plugins.

If plugins contribute commands, they should also implement *getCommandHelp*, which returns a *HelpList*, which is passed a list of Strings. The first word on each help line must match the command name that it describes – e.g. “print <exp> - evaluate an expression”. These are displayed when the user enters the “help” command.

4 Example Plugin Functionality

Hopefully, it is clear how a simple plugin could work, responding to IDE events or providing Commands for the console that add functionality. But this section contains a few examples in more detail, to illustrate what is possible using plugins.

4.1 Extra Type Checking

The “TC” plugin provides type checking for specifications. But it would be possible to add more sophisticated checks via a plugin:

- The plugin has a new name, perhaps “TCX”
- It is likely to want to register for the same events as the “TC” plugin in its *init* method. So it is invoked at each stage of the checking process. The *CheckSyntaxEvent* could be ignored, assuming it is not offering new syntax errors. Most of the work would be done via the *CheckTypeEvent*, or perhaps the *CheckCompleteEvent*.
- The “TC” plugin will have created a TC tree, which can be obtained from the plugin, via the *PluginRegistry*. The *getTC* method will return a *TCModuleList* or a *TCClassList* for the specification. Note that this means you might want SL and PP variants of your plugin, with a factory method (see 3.3).
- The new plugin can process the TC tree and return any extra errors via the return value to the *handleEvent* that received the *CheckTypeEvent*.

4.2 Additional Code Lenses

The built-in plugins contribute the “Launch | Debug” lenses that you see in the Editor view of Visual Studio Code. But extra lenses can be supplied by a plugin, either at the same locations or totally separate locations in a file.

- Note that lenses are very Client specific, because they send back abstract “commands” that the Client should execute, but the meaning of the commands is not defined by LSP.
- Any user plugin can implement the *getCodeLenses* method to provide lenses.
- There is a *CodeLens* abstract class that has helper methods for creating lens responses. This also includes methods to allow you to check the Client type.
- You are likely to want to locate your lenses at a position in the file that corresponds to the location of a VDM construct – like a definition or expression type. To find these locations, you are likely to want to get the parsed/checked tree from the “AST” or “TC” plugins, via the *PluginRegistry*.

4.3 Additional or Replacement Commands

A plugin can contribute executable commands by implementing the *getCommand* and *getCommandHelp* methods. These will then appear in addition to the standard commands in the console when you enter “help”.

- *getCommand* is passed the line as typed by the user. You are responsible for parsing it and deciding whether you recognise the command.
- The *Command* can write to the console (stdout or stderr) via its return value. See the *ExampleCommand* class in the example plugin.
- Remember to implement *getCommandHelp* as well. The help strings passed to the *HelpList* constructor must start with the command name.
- If you give a command the same name as an existing command, yours will replace the earlier

one. This effectively means you can disable a built-in command via a plugin, or you can augment one by calling the original from your own code – to get the original, call the *getCommand* of the plugin that provides it.

- Command objects returned are stateful, so it is best to construct a new Command for every call of *getCommand*, rather than caching objects for reuse.
- If you override the *notWhileRunning* method, and return false, you will be allowed to run your command even while the system is evaluating an expression. Use this with care!
- If you implement the *InitRunnable* interface, you will be able to use your command in the “command” argument in “launch.json”. The interface requires you to implement some methods that will be called in this case, and which allow the result to be formatted as you require.
- If you implement the *ScriptRunnable* interface, you will be able to include your command in scripts (ie. so that the “script” command can execute them). In this case you have to provide a *scriptRun* method as well as the regular *run* method.

Note that when you are stopped at a breakpoint, the console is very restricted in what it can execute, and user-plugin commands are not available.

4.4 Adding New Capabilities

You can turn on LSP/DAP server capabilities in a plugin and then respond to the IDE messages they provoke.

- The *setServerCapabilities* method in your plugin can add or amend the settings added by previous plugins (typically, built-in plugins).
- If the plugin enables (say) “codeActionProvider”, and the Client also supports this, then the IDE will start sending “textDocument/codeAction” requests. These are not handled by the built-in plugins, so an *UnknownMethodEvent* will be published.
- If your plugin registers to handle *UnknownMethodEvents*, it will receive these notifications and (after checking that the method is “textDocument/codeAction”!) can proceed to implement the response to those messages.
- The IDE should then do whatever you requested (usually refactorings, in this case).

5 Future Work

The language server is constantly being worked on, though we try to keep versions stable! The following are areas that need more work, or are in development (as of 4.5.0-SNAPSHOT):

- Better DAP protocol integration via events.
- Plugins that include Client side components too.

A. Events

The following events are available via the *EventHub*:

Event Class	When?
ChangeFileEvent	
CheckCompleteEvent	
CheckPrepareEvent	
CheckSyntaxEvent	
CheckTypeEvent	
CloseFileEvent	
DAPConfigDoneEvent	
DAPDisconnectEvent	
DAPEvaluateEvent	
DAPInitializeEvent	
DAPLaunchEvent	
DAPTerminateEvent	
InitializedEvent	
InitializeEvent	
OpenFileEvent	
SaveFileEvent	
ShutdownEvent	
UnknownCommandEvent	
UnknownMethodEvent	
UnknownTranslationEvent	